

LM efficiency trade-off

Dev PPL versus throughput. Lower PPL and higher tokens/s are preferable; point size follows parameter count.

