

LM results on Penn Treebank

Perplexity comparison. Lower is better. LoRA uses frozen pretrained GPT2 plus manual adapters.

