

LM/partA train loss across all mandatory experiments

All core scratch GPT2 runs. This dense view is useful for diagnostics rather than the final one-page report.

