

NLU/partB loss curves: BERT vs GPT2 core

Total multitask loss. Solid=train, dashed=dev. BERT converges to a much lower dev loss.

